

```

import numpy as np
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

class LogisticRegression:
    def __init__(self, lr = 0.01, batch_size = 16, iters = 10):
        #initialize the learning rate, batch size, iterations, and weights
        self.lr = lr
        self.batch_size = batch_size
        self.iters = iters
        self.weights = None

    def sigmoid(self, x):
        #create function from sigmoid
        return 1/(1+np.exp(-x))

    def loss(self, y, y_pred):
        #calculate loss of a data sample
        return (-1/y.shape[0])*np.sum(y*np.log(y_pred+1e-8) + (1-y)*np.log(1-y_pred+1e-8))

    def grad(self, X, y, y_pred):
        #calculate gradient of data sample with identity feature mapping
        return (-1/y.shape[0])*(X.T).dot(y-y_pred)

    def predict(self, X):
        #predict results based on weights
        return self.sigmoid(np.dot(X, self.weights))

    def train(self, X, y):
        samples, features = X.shape
        self.weights = np.random.normal(size = features)

        for i in range(self.iters):
            #shuffle indices
            inds = np.arange(samples)
            np.random.shuffle(inds)
            X_alt = X[inds]
            y_alt = y[inds]

            for j in range(samples//self.batch_size):
                #create batches
                X_batch = X_alt[j*self.batch_size:(j+1)*self.batch_size]
                y_batch = y_alt[j*self.batch_size:(j+1)*self.batch_size]

                #predict results of batch using current weights
                y_pred = self.predict(X_batch)

                #update gradient
                grad = self.grad(X_batch, y_batch, y_pred)
                self.weights -= self.lr*grad

            #calculate loss
            y_pred = self.predict(X)
            loss = self.loss(y, y_pred)

        #print("Final Loss: ", loss, " for ", self.lr, " ", self.batch_size, " ", self.iters)

data = load_breast_cancer()
X = data['data']
y = data['target']
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size = 0.2)
X_val, X_test, y_val, y_test = train_test_split(X_val, y_val, test_size = 0.5)

print("X train size: ", X_train.shape)
print("y train size: ", y_train.shape)
print("X validation size: ", X_val.shape)
print("y validation size: ", y_val.shape)
print("X test size: ", X_test.shape)
print("y test size: ", y_test.shape)

```

```

X train size: (455, 30)
y train size: (455,)
X validation size: (57, 30)
y validation size: (57,)
X test size: (57, 30)
y test size: (57,)

for lr in [1, 1e-2, 1e-3, 1e-4, 1e-5]:
    for batch_size in [4, 8, 16, 32, 64]:
        error = 0
        for i in range(10):
            LR = LogisticRegression(lr = lr, batch_size = batch_size)
            LR.train(X_train, y_train)

            y_pred = LR.predict(X_val)
            y_pred = np.where(y_pred > 0.5, 1, 0)
            error += np.sum(y_pred != y_val)/y_val.shape[0]

        print("Error: ", error/10, " for ", lr, " ", batch_size)

<ipython-input-109-472a90248188>:9: RuntimeWarning: overflow encountered in exp
    return 1/(1+np.exp(-x))
Error: 0.2192982456140351 for 1 4
Error: 0.2245614035087719 for 1 8
Error: 0.23684210526315788 for 1 16
Error: 0.28771929824561404 for 1 32
Error: 0.41052631578947363 for 1 64
Error: 0.19649122807017544 for 0.01 4
Error: 0.15438596491228068 for 0.01 8
Error: 0.30877192982456136 for 0.01 16
Error: 0.3280701754385965 for 0.01 32
Error: 0.43508771929824563 for 0.01 64
Error: 0.21052631578947367 for 0.001 4
Error: 0.18771929824561404 for 0.001 8
Error: 0.20526315789473681 for 0.001 16
Error: 0.34385964912280703 for 0.001 32
Error: 0.36140350877192984 for 0.001 64
Error: 0.18596491228070172 for 0.0001 4
Error: 0.25614035087719295 for 0.0001 8
Error: 0.2649122807017544 for 0.0001 16
Error: 0.4736842105263158 for 0.0001 32
Error: 0.44035087719298244 for 0.0001 64
Error: 0.3421052631578948 for 1e-05 4
Error: 0.33508771929824566 for 1e-05 8
Error: 0.41052631578947363 for 1e-05 16
Error: 0.4491228070175438 for 1e-05 32
Error: 0.5333333333333333 for 1e-05 64

```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
#choose lr of 0.01 and batch size of 8 because it gives lowest error rate
```

```
LR = LogisticRegression(lr = 0.01, batch_size = 8)
loss = LR.train(X_train, y_train)
```

```
y_pred = LR.predict(X_test)
y_pred = np.where(y_pred > 0.5, 1, 0)
```

```
print("Accuracy: ", accuracy_score(y_test, y_pred))
print("Precision: ", precision_score(y_test, y_pred))
print("Recall: ", recall_score(y_test, y_pred))
print("F1 Score: ", f1_score(y_test, y_pred))
```

```

Accuracy: 0.8421052631578947
Precision: 0.7857142857142857
Recall: 1.0
F1 Score: 0.88
<ipython-input-109-472a90248188>:9: RuntimeWarning: overflow encountered in exp
    return 1/(1+np.exp(-x))

```

The accuracy of the model is 0.84, which means that it will distinguish whether a patient has breast cancer or not 84% of the time. The precision is 0.78, which means there are a few false positive results. This will predict that some people have breast cancer when they do not actually have breast cancer. The recall score is 1, which means there are no false negatives, so no one will be told they don't have breast cancer but they do. A high recall score is critical in the case of breast cancer since it can be deadly if breast cancer is missed. F1 score combines the precision and recall scores. In this case, the F1 score is 0.88.

Overall, the most important score is recall for this model. Since the recall is 1 for this model, it performs quite well.